

Mini-Project MA-ICR

A shared encrypted network file system

Titus Abele

MSE Computer Science
HES-SO Master
Lausanne, Switzerland
May 6, 2024

Contents

1	Architecture	4
1.1	Web Server	4
1.2	Storage	5

Listings

1	Ubuntu root folder	2
2	Trying to create a file in the root folder without write permissions	3

List of Figures

1	General architecture of a web server with file saving and sharing capabilities	4
2	General interaction graph	4
3	A basic filesystem	5

Introduction

Before jumping into the theory behind a shared encrypted network file system, it is important to lay a foundation about what a file system is, why it requires sharing capabilities and why security is paramount. A file system (FS) manages and provides access for resources. Generally, resources are composed of folders containing files or other folders. This way, an organizational hierarchy of resources can be established.

```
[drwxr-xr-x 4.0K] .
|--- [lrwxrwxrwx 7] bin -> usr/bin
|--- [drwxr-xr-x 4.0K] boot
|--- [drwxr-xr-x 3.5K] dev
|--- [drwxr-xr-x 4.0K] etc
|--- [drwxr-xr-x 4.0K] home
|--- [rwxrwxrwx 2.0M] init
|--- [lrwxrwxrwx 7] lib -> usr/lib
|--- [lrwxrwxrwx 9] lib32 -> usr/lib32
|--- [lrwxrwxrwx 9] lib64 -> usr/lib64
|--- [lrwxrwxrwx 10] libx32 -> usr/libx32
|--- [drwx----- 16K] lost+found
|--- [drwxr-xr-x 4.0K] media
|--- [drwxr-xr-x 4.0K] mnt
|--- [drwxr-xr-x 4.0K] opt
|--- [dr-xr-xr-x 0] proc
|--- [drwx----- 4.0K] root
|--- [drwxr-xr-x 680] run
|--- [lrwxrwxrwx 8] sbin -> usr/sbin
|--- [drwxr-xr-x 4.0K] snap
|--- [drwxr-xr-x 4.0K] srv
|--- [dr-xr-xr-x 0] sys
|--- [drwxrwxrwt 12K] tmp
|--- [drwxr-xr-x 4.0K] usr
|--- [drwxr-xr-x 4.0K] var

23 directories, 1 file
```

Listing 1: Ubuntu root folder

Access should be handled so that some roles can access some resources, this is generally done by using dedicated roles such as administrators, owners, guests or even users. An example can be seen in Listing 1, it shows all folders and files contained within the root directory of a machine running Ubuntu, a popular Linux distribution. The root directory is the top-level directory in the file system's hierarchy. Left of the resource names (**boot**, **dev**, **etc**...) we can find their relative permissions. For instance, the **boot** directory is marked as `[drwxr-xr-x 4.0K] boot` which means:

- The **d** in first position indicates the nature of the resource, in this case a directory.
- The pattern **rw****xr****-xr****-x** translates to the owner having read, write, and execute permissions, while the group and others have read and execute permissions only.

Each triplet (**rw****x**) relates to a specific permission class. In a Unix-like file system such as the one depicted in Listing 1, these classes are defined in order as "Owner-Group-Others". This means that for the **boot** directory:

- The Owner has read (**r**), write (**w**) and execute (**x**) permissions.
- The Group that is associated with the directory and all other users only have read and execute permissions.

This way any access to the directory is dynamically limited and permissions can be revoked, approved and modified very easily.

```
seirios@T16:/$ touch foo.txt
touch: cannot touch 'foo.txt': Permission denied
```

Listing 2: Trying to create a file in the root folder without write permissions

The reason such strategies are put in place is not only related to security. The root directory, such as the one from the Linux distribution, contains all the resources necessary for the proper functioning of the operating system (OS). Modifying these files may cause irreversible configurations that may break the proper flow of the OS and cause failure. The user trying to create a text file using the **touch**[1] command in Listing 2 is being denied by the operating system. The current directory being the root directory, one needs so called root-privileges to create, modify or delete any resources¹.

To further emphasize the importance of the mechanic surrounding permissions, we must talk about security. A machine such as the one depicted above, may be of use to multiple actors. Therefore, systems must be in place to regulate access across user sessions and resources. Alice may not want Bob to read, edit or delete some of her files or even worse create some in her name. This means that a file system must also be able to obfuscate files from eavesdroppers and make them unavailable, unreadable and uneditable to malicious attackers. This means that in order to consult her own files, Alice must be logged in otherwise her files are inaccessible. But there might be a file which Alice wants to share with Bob for a project, in this case she may create a group, add Bob to this group and using the appropriate triplet from earlier, specify that people of that group may read and write to this file.

Project goal

The goal of this project is to design a shared encrypted network file system. The most important aspect of this FS is the manner in which security is guaranteed. In this report, we will outline the general structure of the FS and for each component, specify the necessary cryptographic tools used to provide trust and security.

¹There are many ways of obtaining some of these privileges notably the **sudo** command which will give the user elevated permissions. It is important to note however, that the user trying to **sudo** a command must be part of a **sudoers** group

1 Architecture

1.1 Web Server

A web server is an online machine that serves clients. Its purpose can be manifold. In our project, this web server will store files for users. They may connect to the server using some credentials (in our case a unique username and a password) and have access to their files as well as the files that have been shared with them by other users. In Figure 1 we outlined the general architecture of the server. The server prompts the client for credentials, the client provides them, the server can then verify the credentials and grant access to the storage. In this way, the client can connect from anywhere and on any device, without having to transport the entire storage. This is essentially what is commonly referred to as a cloud storage service.




Figure 1: General architecture of a web server with file saving and sharing capabilities

The standard procedure to interact with the server is outlined in Figure 2. As depicted, there are four components that need to be conceptualized.

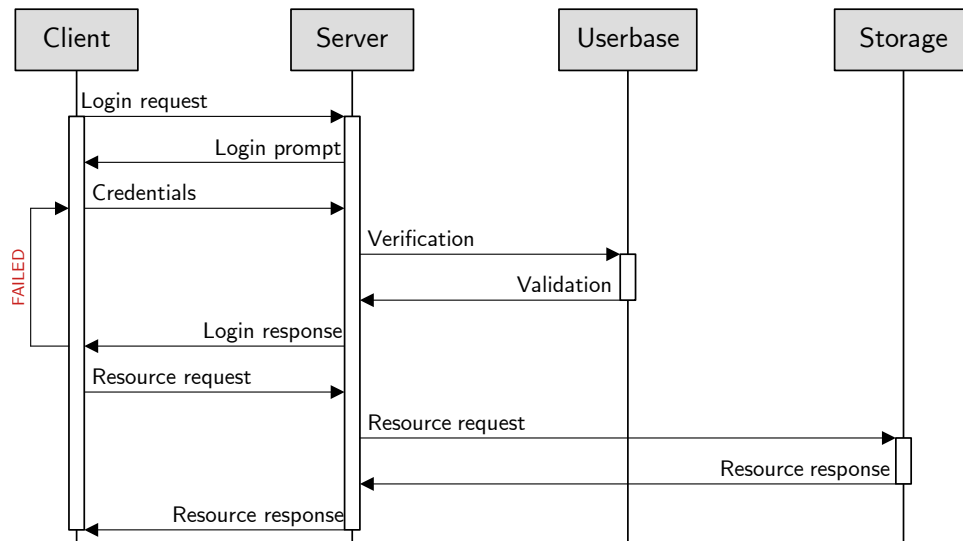


Figure 2: General interaction graph

1.2 Storage




Figure 3: A basic filesystem

The building blocks for a simple file system are as follows:

- Folders are containers for other folders and resources.
- Resources are bundles of information that manifest in the form of files.
- A root folder which is the root node of the hierarchy that illustrates the file system.

In Figure 3 the root folder is the left most folder.

A file can therefore be defined as a unique place in the file system by a specific identifier which indicates the path one must travel to find it. To be able to retrieve the files in the green folder in Figure 3, one must start to go into the second folder in the root folder and then move to the first folder contained within. So if we translate this to a path, we would have something akin to `\root\second_folder\first_folder` which could be defined as the unique identifier for the green folder. We may omit the `\root` indicator because all resources are contained within the `root` directory and also replace the descriptive names with incremental identification numbers (`first_folder` would therefore carry the identifier 0, `second_folder` would carry 1 and so on). The resulting unique identifier for the green folder would therefore be: 10. Similarly, the blue folder is identified by 00. A file could be described in the same way.

References

- [1] The `touch` Command Manual Page - Linux manual page. (n.d.).
<https://man7.org/linux/man-pages/man1/touch.1.html>